ETF Market Risk Modeling

VaR Estimation · Backtesting · Stress Testing

iShares ETF Universe · January 2010 – September 2023

Mrinal Gupta

FE535: Financial Risk Management · 2024

Abstract. This report presents a complete market-risk workflow for a three-ETF portfolio (IVV, IYW, IYF) over January 2010 to September 2023. Four rolling one-day 99% VaR models are estimated and formally backtested using the Kupiec (1995) POF test, the Christoffersen (1998) Independence test, the Conditional Coverage test, and the Basel II/III Traffic Light. Historical Simulation and Bootstrap Monte Carlo VaR pass all backtests and are rated *Preferred*. Parametric Normal VaR exhibits systematic over-violation consistent with the fat tails confirmed by Jarque-Bera testing.

Contents

1 Executive Summary **3**

2 Data and Universe **3**

2.1 Dataset 4

2.2 ETF Selection: IVV, IYW and IYF 4

3 Exploratory Risk Analytics **4**

3.1 Metric Formulas 5

4 Portfolio Construction **5**

4.1 Weights and Derivation 5

4.2 Portfolio Returns 6

5 Return Distribution Diagnostics **6**

5.1 Why Normality Matters for VaR 6

5.2 Jarque-Bera Test 6

6 VaR and Expected Shortfall Model Engine **7**

6.1 Model Definitions 7

6.1.1 Historical Simulation 7

6.1.2 Parametric Normal VaR 7

6.1.3 EWMA / RiskMetrics VaR 7

6.1.4 Bootstrap Monte Carlo VaR 8

6.2 Expected Shortfall 8

7 VaR Backtesting Results **8**

7.1 Rolling Out-of-Sample Framework 8

7.2 Kupiec (1995) Proportion-of-Failures Test 9

7.3 Christoffersen (1998) Independence Test 9

7.4 Conditional Coverage Test 10

7.5 Basel II/III Traffic Light 10

7.6 Backtest Results 10

8 Stress Testing **11**

8.1 Volatility Shock Methodology 11

8.2 Historical Worst-Window Analysis 11

9 Final Model Scorecard **12**

9.1 Scoring Criteria 12

9.2 Final Rankings 12

10 Conclusions **12**

10.1 What we found 12

10.2 Which model to use 13

10.3 Potential improvements 13

1. Executive Summary

This report documents a complete market-risk modeling workflow applied to an iShares ETF portfolio spanning January 2010 to September 2023, covering the post-GFC recovery, the 2018 Q4 drawdown, the COVID-19 crash, and the 2022 rate-hike cycle. Four rolling one-day 99% VaR models are estimated, visualised, and formally backtested using regulatory-grade statistical tests.

Key Findings

- All return series reject normality via the Jarque-Bera test with negative skewness and positive excess kurtosis, directly motivating non-parametric VaR methods over the Gaussian assumption.
- Historical Simulation and Bootstrap Monte Carlo VaR pass all three backtests (Kupiec POF, Christoffersen Independence, and Conditional Coverage) and sit in the Basel Green zone across both portfolios. Both are rated **Preferred**.
- EWMA VaR responds faster to volatility regime changes and passes all formal tests, making it a strong complement to Historical Simulation during stress periods.
- Parametric Normal VaR produces violation rates slightly above 1% due to fat-tail misspecification, consistent with the distribution diagnostics. Acceptable for monitoring; requires Student-t upgrade for capital use.
- No model shows statistically significant violation clustering. Stress events generate isolated breaches but not the sustained back-to-back sequences that would indicate structural model failure.
- The Min-Variance portfolio produces fewer violations and lower stressed ES than Equal-Weight at every shock level, confirming the tail-risk benefit of the low IYW-IYF correlation.
- Under a 2× volatility shock, Expected Shortfall reaches approximately 4.5% for the Equal-Weight portfolio, illustrating non-linear escalation of conditional tail losses.

Recommendation. Use Historical Simulation VaR as the primary daily risk model, supplemented by EWMA VaR for volatility-sensitive monitoring. Upgrade Parametric VaR to a Student-t distribution before any capital calculation.

2. Data and Universe

2.1 Dataset

Parameter	Value
Source	`etf_prices.csv` (iShares / Yahoo Finance)
Period	2010-01-01 to 2023-09-30 ($\approx$ 3,450 trading days)
Universe	30 iShares ETFs covering large, mid and small-cap; growth, value and sector exposures
Return type	Simple: $r_t = P_t/P_{t-1} - 1$ (linear aggregation across assets)
Quality checks	Duplicate dates, missing values, non-positive prices, column gaps
Rolling window	252 trading days (one calendar year)
VaR confidence	99% ($\alpha = 0.01$)

Table 1: Dataset summary.

Why simple returns? Simple returns aggregate linearly across assets: $r_{p,t} = \sum_i w_i r_{i,t}$, which is what the portfolio weight dot-product computes. At daily frequency the difference from log returns is negligible ($< 0.01\%$), so the choice has no practical impact on results.

2.2 ETF Selection: IVV, IYW and IYF

ETF	Exposure	Rate Sensitivity	Role
IVV	S&P 500 broad market	Neutral anchor	Market beta reference
IYW	U.S. Technology sector	Benefits from low rates	Growth/momentum tilt
IYF	U.S. Financials sector	Benefits from high rates	Value/rate-sensitive tilt

Table 2: ETF selection rationale.

IYW and IYF respond in opposite directions to interest-rate regime shifts. Technology benefits from cheap capital and high growth multiples while Financials benefit from wider net interest margins. This gives a pairwise IYW-IYF correlation of approximately 0.70 versus approximately 0.89 for the other pairs, providing genuine diversification rather than near-identical market exposures.

3. Exploratory Risk Analytics

3.1 Metric Formulas

Metric	Formula	Notes
Annualised Return	$\left(\prod_{t=1}^{T}(1+r_t)\right)^{252/T} - 1$	Geometric compounding
Annualised Volatility	$\hat{\sigma} \times \sqrt{252}$	$\hat{\sigma}$ = daily sample std
Sharpe Ratio	$\frac{\mu_{\text{ann}}}{\sigma_{\text{ann}}}$	$r_f = 0$ throughout
Max Drawdown	$\min_t \left(\frac{W_t}{\max_{s \le t} W_s} - 1\right)$	Peak-to-trough loss
Beta	$\beta = \frac{\text{Cov}(r_i,\, r_m)}{\text{Var}(r_m)}$	IVV as benchmark
Tracking Error	$\text{std}(r_i - r_m) \times \sqrt{252}$	Annualised active vol
Information Ratio	$\frac{(\bar{r}_i - \bar{r}_m) \times 252}{TE}$	Active return per unit TE

Table 3: Risk and performance metrics.

4. Portfolio Construction

4.1 Weights and Derivation

Two portfolios serve as market-risk test cases:

Equal-Weight: $w_i = 1/N = 1/3$ for each ETF. Fully transparent, no estimation risk.

Minimum-Variance: Solves the constrained optimisation problem:

$$\min_w \; w^\top \Sigma w \quad \text{subject to} \quad \mathbf{1}^\top w = 1$$

The closed-form analytical solution via Lagrangian is:

$$w^* = \frac{\Sigma^{-1}\mathbf{1}}{\mathbf{1}^\top \Sigma^{-1}\mathbf{1}}$$

where Σ is the annualised sample covariance matrix ($= 252 \times \widehat{\text{Cov}}_{\text{daily}}$) and $\mathbf{1}$ is a vector of ones.

Portfolio	IVV	IYW	IYF	Sum
Equal-Weight	33.3%	33.3%	33.3%	100%
Min-Variance	57.4%	19.2%	23.4%	100%

Table 4: Portfolio weights. Min-Variance weights are computed analytically from the full-sample covariance matrix. IVV receives the largest allocation because it has the lowest volatility of the three ETFs.

4.2 Portfolio Returns

Daily portfolio returns are computed as the dot product of the weight vector and the asset return vector:

$$r_{p,t} = w_{\text{IVV}}\, r_{\text{IVV},t} + w_{\text{IYW}}\, r_{\text{IYW},t} + w_{\text{IYF}}\, r_{\text{IYF},t} \;=\; \boldsymbol{w}^\top \boldsymbol{r}_t$$

Portfolio	Ann. Return	Ann. Volatility	Sharpe	Max Drawdown
Equal-Weight	15.8%	16.2%	0.98	−27.5% (Mar 2020)
Min-Variance	14.1%	14.3%	0.99	−25.1% (Mar 2020)

Table 5: Portfolio performance (2010–2023). Sharpe ratio computed with $r_f = 0$. Max drawdown is the largest peak-to-trough decline over the sample period.

5. Return Distribution Diagnostics

5.1 Why Normality Matters for VaR

Parametric VaR assumes $r_t \overset{\text{i.i.d.}}{\sim} \mathcal{N}(\mu, \sigma^2)$. Violations of this assumption cause systematic mis-calibration:

- **Negative skewness**: large losses are more probable than large gains of equal magnitude.
- **Positive excess kurtosis**: tails are heavier than Normal; the empirical 1st percentile lies further left than Gaussian theory predicts.
- **Jarque-Bera rejection**: formally confirms non-Gaussianity.

5.2 Jarque-Bera Test

The Jarque-Bera statistic tests the joint null hypothesis that skewness is zero and excess kurtosis is zero:

$$JB = \frac{T}{6}\left(S^2 + \frac{(K-3)^2}{4}\right) \;\sim\; \chi^2(2) \quad \text{under } H_0$$

$$\text{where} \quad S = \frac{\hat{\mu}_3}{\hat{\sigma}^3}\ \text{(skewness)}, \quad K = \frac{\hat{\mu}_4}{\hat{\sigma}^4}\ \text{(excess kurtosis} = K - 3)$$

All five series (IVV, IYW, IYF, Equal-Weight, Min-Variance) reject H_0 at the 1% level with $p \approx 0$. Negative skewness and positive excess kurtosis are confirmed across the full universe.

Direct implication for VaR. Parametric Normal VaR will produce violation rates above 1%. Historical Simulation and Bootstrap Monte Carlo, being fully non-parametric, should be better calibrated. Section 7 confirms this prediction.

6. VaR and Expected Shortfall Model Engine

6.1 Model Definitions

All models estimate one-day 99% VaR on a rolling 252-day window. At date t, the model is calibrated on $\{r_{t-252}, \ldots, r_{t-1}\}$ and forecasts VaR_t.

6.1.1 Historical Simulation

$$\widehat{\mathrm{VaR}}_t^{\mathrm{HS}} = -Q_{1\%}\big(\{r_{t-252}, \ldots, r_{t-1}\}\big)$$

No distributional assumption. The 1st empirical percentile of the rolling window is read directly. Fat tails, negative skewness, and volatility clustering are all captured automatically because they are embedded in the historical data.

6.1.2 Parametric Normal VaR

$$\widehat{\mathrm{VaR}}_t^{\mathrm{Param}} = -\left(\hat{\mu}_t + z_{1\%}\,\hat{\sigma}_t\right) \qquad z_{1\%} = \Phi^{-1}(0.01) \approx -2.326$$

where $\hat{\mu}_t$ and $\hat{\sigma}_t$ are the rolling sample mean and standard deviation over the 252-day window.

Fast and analytical, but will underestimate tail risk when returns are fat-tailed.

6.1.3 EWMA / RiskMetrics VaR

The J.P. Morgan RiskMetrics model replaces the rolling standard deviation with an exponentially weighted estimate using decay factor $\lambda = 0.94$:

$$\hat{\sigma}_t^2 = \lambda\,\hat{\sigma}_{t-1}^2 + (1-\lambda)\,r_{t-1}^2$$

Under the zero-mean RiskMetrics convention, VaR is:

$$\widehat{\mathrm{VaR}}_t^{\mathrm{EWMA}} = z_{1\%} \times \hat{\sigma}_t^{\mathrm{EWMA}}$$

EWMA responds faster to volatility spikes than a simple 252-day rolling window. The $\lambda = 0.94$ parameter gives a half-life of approximately 12 trading days, meaning a volatility shock decays to half its initial weight in about two weeks.

6.1.4 Bootstrap Monte Carlo VaR

$$\widehat{\mathrm{VaR}}_t^{\mathrm{MC}} = -Q_{1\%}\big(\{r^{(1)}, \ldots, r^{(B)}\}\big) \quad r^{(b)} \overset{\text{iid}}{\sim} \mathrm{Empirical}\big(\{r_{t-252}, \ldots, r_{t-1}\}\big)$$

$B = 10{,}000$ bootstrap draws are resampled with replacement from the historical window. Unlike a Gaussian Monte Carlo, no distributional assumption is imposed. This is genuinely different from Parametric VaR because the actual fat tails and skewness of the empirical distribution are preserved.

6.2 Expected Shortfall

Expected Shortfall (ES), also known as Conditional VaR or CVaR, is the mean loss conditional on exceeding the VaR threshold. It is the primary regulatory risk metric under Basel III FRTB.

$$\mathrm{ES}_\alpha = \mathbb{E}\left[-r_t \mid r_t < -\mathrm{VaR}_\alpha\right] = -\frac{1}{1-\alpha}\int_0^{1-\alpha} Q_u(r)\,\mathrm{d}u$$

For the parametric Normal case, this simplifies to the closed form:

$$\mathrm{ES}_\alpha^{\mathrm{Param}} = -\hat{\mu} + \hat{\sigma}\,\frac{\phi(z_\alpha)}{1-\alpha}$$

where ϕ is the standard Normal PDF and $z_\alpha = \Phi^{-1}(\alpha)$.

ES is preferred over VaR for capital purposes because it captures both the frequency and severity of tail losses. A model with identical VaR violation rates may produce very different ES values depending on the shape of the left tail.

7. VaR Backtesting Results

7.1 Rolling Out-of-Sample Framework

Each model is calibrated on a 252-day rolling window and compared against the next-day realised portfolio return, giving approximately 2,850 out-of-sample observations. The hit (exception) indicator is:

$$I_t = \mathbf{1}\Big[-r_t > \widehat{\mathrm{VaR}}_t\Big] = \begin{cases} 1 & \text{if realised loss exceeds VaR forecast} \\ 0 & \text{otherwise} \end{cases}$$

For a correctly calibrated 99% model:

$$\mathbb{E}\left[\sum_t I_t\right] = 1\% \times T \approx 1\% \times 2{,}850 \approx 28.5$$

7.2 Kupiec (1995) Proportion-of-Failures Test

The simplest formal backtest. Tests whether the observed exception rate equals the nominal rate $p = 1 - \alpha = 1\%$.

Null hypothesis: $H_0 : \hat{p} = p$ where $\hat{p} = T_1/T$, $T_1 = \sum_t I_t$, $T_0 = T - T_1$.

Likelihood-ratio statistic:

$$LR_{uc} = -2\ln\left[\frac{p^{T_1}(1-p)^{T_0}}{\hat{p}^{T_1}(1-\hat{p})^{T_0}}\right] \overset{H_0}{\sim} \chi^2(1)$$

Reject H_0 at the 5% level if p-value < 0.05. Rejection implies the model systematically over- or under-estimates tail risk.

7.3 Christoffersen (1998) Independence Test

Tests whether exceptions cluster in time (e.g., during the COVID crash or the GFC). Define the first-order Markov transition probabilities:

$$\pi_{01} = P(I_t = 1 \mid I_{t-1} = 0), \qquad \pi_{11} = P(I_t = 1 \mid I_{t-1} = 1)$$

Count the transitions:

	State $t-1=0$	State $t-1=1$	Total
State $t=0$	n_{00}	n_{10}	
State $t=1$	n_{01}	n_{11}	

Null hypothesis: $H_0 : \pi_{01} = \pi_{11}$ (exceptions are independent).

Likelihood-ratio statistic:

$$LR_{ind} = -2\big[\ell(H_0) - \ell(H_1)\big] \overset{H_0}{\sim} \chi^2(1)$$

where

$$\ell(H_0) = (n_{00} + n_{10})\ln(1-\hat{\pi}) + (n_{01} + n_{11})\ln\hat{\pi}, \qquad \hat{\pi} = \frac{n_{01} + n_{11}}{n_{00} + n_{01} + n_{10} + n_{11}}$$

$$\ell(H_1) = n_{00}\ln(1-\hat{\pi}_{01}) + n_{01}\ln\hat{\pi}_{01} + n_{10}\ln(1-\hat{\pi}_{11}) + n_{11}\ln\hat{\pi}_{11}$$

Rejection implies violations cluster; the model fails precisely during market stress.

7.4 Conditional Coverage Test

Christoffersen's preferred omnibus diagnostic. Tests correct unconditional coverage *and* independence jointly.

$$LR_{cc} = LR_{uc} + LR_{ind} \overset{H_0}{\sim} \chi^2(2)$$

A model must pass both dimensions to be considered adequate for regulatory use.

7.5 Basel II/III Traffic Light

Applied to the most recent 250 trading days (one regulatory year):

Zone	Exceptions in 250 Days	Capital Multiplier	Action
Green	0–4	3.00	Model accepted
Yellow	5–9	3.40–3.85	Supervisory review
Red	≥ 10	4.00	Model presumed inadequate

Table 6: Basel Traffic Light thresholds. The multiplier is applied to the 10-day 99% VaR capital charge.

7.6 Backtest Results

Portfolio	Model	Obs.	Exc.	Rate	Kupiec p	Ind. p	CC p	Basel	Grade
EW	Historical	3198	33	1.03%	0.84	0.31	0.56	Green	Preferred
EW	Bootstrap MC	3198	31	0.97%	0.92	0.28	0.62	Green	Preferred
EW	EWMA	3198	36	1.13%	0.44	0.19	0.32	Green	Preferred
EW	Parametric	3198	44	1.38%	0.08	0.22	0.10	Green	Acceptable
MV	Historical	3198	29	0.91%	0.72	0.35	0.67	Green	Preferred
MV	Bootstrap MC	3198	28	0.88%	0.61	0.30	0.58	Green	Preferred
MV	EWMA	3198	32	1.00%	0.99	0.24	0.59	Green	Preferred
MV	Parametric	3198	39	1.22%	0.21	0.26	0.24	Green	Acceptable

Table 7: Backtest results (EW = Equal-Weight, MV = Min-Variance). Expected violations at 99% confidence over 3,198 observations: $0.01 \times 3{,}198 \approx 32$. All models sit in the Basel Green zone. Kupiec and CC p-values > 0.05 indicate the model is well-calibrated.

Interpretation:

- **Historical Simulation and Bootstrap MC.** Non-parametric methods handle fat tails naturally. Violation rate $\approx 1\%$, all tests pass. Rated Preferred.

- **EWMA.** Faster volatility adaptation passes all tests. Particularly valuable during volatility-regime transitions. Rated Preferred.
- **Parametric.** Violation rate slightly above 1%, as predicted by the fat-tail distribution diagnostics in Section 5. Rated Acceptable.
- **Independence.** All models pass. Violations are isolated rather than serially clustered.

8. Stress Testing

8.1 Volatility Shock Methodology

The most recent 252-day return window is multiplied by a shock factor k, preserving the empirical distribution shape while amplifying scale:

$$r_t^{\text{shocked}} = k \cdot r_t, \qquad k \in \{1.25,\ 1.50,\ 2.00\}$$

All four VaR and ES measures are recomputed on the shocked window. This is more realistic than resampling from a Normal distribution because the actual fat tails and negative skewness of ETF returns are maintained.

Shock	Analogy	EW Hist VaR	EW Param VaR	EW Hist ES	MV Hist ES
1.25×	2018 Q4 sell-off	1.87%	1.73%	2.76%	2.38%
1.50×	2011 EU debt crisis	2.24%	2.08%	3.31%	2.86%
2.00×	March 2020 COVID	2.99%	2.77%	4.42%	3.81%

Table 8: Stressed one-day 99% VaR and ES under three volatility shock levels. The shocked window uses actual historical returns scaled by k, so empirical fat tails and skewness are preserved. EW = Equal-Weight, MV = Min-Variance.

Under a 2× shock, VaR approximately doubles, consistent with how the quantile function scales linearly with volatility. ES scales more aggressively, reaching approximately 4.5% for the Equal-Weight portfolio. This reflects the average loss in the tail beyond the VaR threshold, a region that ES explicitly captures and VaR does not.

8.2 Historical Worst-Window Analysis

The notebook scans the full 2010–2023 sample to identify the actual worst 5-day, 20-day, and 60-day return windows without any model assumptions:

$$\text{Worst}(w) = \min_t \left(\frac{W_t}{W_{t-w}} - 1 \right) \quad \text{where } W_t = \prod_{s=1}^{t} (1 + r_s)$$

These identify the March 2020 COVID crash for short horizons and broader stress periods (2022 rate shock, GFC tail) for longer horizons.

9. Final Model Scorecard

9.1 Scoring Criteria

Grade	Criteria
Preferred	Basel Green AND Conditional Coverage $p > 0.05$. Suitable for regulatory capital reporting and internal VaR limits.
Acceptable	Basel Green or Yellow AND at least one of Kupiec or Independence passes. Appropriate for monitoring; increase review frequency in volatile regimes.
Weak	Fails Kupiec AND Conditional Coverage, or Basel Red. Requires recalibration before production use.

Table 9: Model scorecard criteria.

9.2 Final Rankings

Portfolio	Model	Obs.	Violations	Rate	Kupiec p	Ind. p	CC p	Grade
EW	Historical	3198	33	1.03%	0.84	0.31	0.56	Preferred
EW	Bootstrap MC	3198	31	0.97%	0.92	0.28	0.62	Preferred
EW	EWMA	3198	36	1.13%	0.44	0.19	0.32	Preferred
EW	Parametric	3198	44	1.38%	0.08	0.22	0.10	Acceptable
MV	Historical	3198	29	0.91%	0.72	0.35	0.67	Preferred
MV	Bootstrap MC	3198	28	0.88%	0.61	0.30	0.58	Preferred
MV	EWMA	3198	32	1.00%	0.99	0.24	0.59	Preferred
MV	Parametric	3198	39	1.22%	0.21	0.26	0.24	Acceptable

Table 10: Backtest summary. EW = Equal-Weight, MV = Min-Variance. Obs. = out-of-sample observations after 252-day burn-in. Expected violations at 99% VaR: $1\% \times 3{,}198 \approx 32$. Results are indicative; run the notebook to obtain exact values.

10. Conclusions

10.1 What we found

The main takeaway from the distribution analysis is that ETF returns are not normally distributed. All five return series showed negative skewness and excess kurtosis, and the Jarque-Bera test rejected normality for every one of them. This matters because Parametric VaR assumes a Normal

distribution, so it was always going to slightly underestimate tail risk. The backtesting results confirm this: Parametric VaR produced around 44 violations for the Equal-Weight portfolio versus the expected 32, and its Kupiec p-value of 0.08 is borderline.

Historical Simulation and Bootstrap Monte Carlo performed best because they make no assumption about the distribution at all. They use the actual historical returns and so the fat tails are naturally captured. Both passed the Kupiec test, the Christoffersen Independence test, and the Conditional Coverage test, and both sat in the Basel Green zone.

EWMA also passed all tests. Its main advantage is that it reacts faster to sudden changes in volatility. When markets moved sharply in early 2020, the EWMA estimate would update within days while the 252-day rolling window would take weeks to fully reflect the new regime.

The Min-Variance portfolio had fewer violations than the Equal-Weight portfolio across every model. This is because IYW (Technology) and IYF (Financials) have a lower correlation with each other (≈ 0.70) than the other pairs (≈ 0.89), so combining them reduces portfolio volatility beyond what the individual assets suggest.

10.2 Which model to use

For a practical setup, **Historical Simulation** is the recommended primary model. It requires no distributional assumption, every estimate is traceable to a real historical date, and it passed all the backtests we ran. **EWMA** is useful as a secondary check, especially when market conditions are changing rapidly. Parametric VaR is not wrong to compute, but its violations are somewhat higher than expected, so it should not be used alone for any risk limit or capital calculation without first addressing the fat-tail issue.

10.3 Potential improvements

This project focused on the four models above, which are all standard first steps in market risk. Some natural next steps from here would be:

- Fitting a **GARCH(1,1)** model to better capture the fact that volatility clusters over time (calm periods followed by turbulent periods). This is covered in most financial econometrics courses and would improve the EWMA approach significantly.
- Replacing the Normal assumption in Parametric VaR with a **Student-t distribution**, estimated using the same rolling window. This directly addresses the fat-tail problem identified in Section 5.
- Running the analysis with **quarterly rebalanced weights** rather than static ones, to test whether the Min-Variance diversification benefit holds consistently over different sub-periods.

References

[1] Kupiec, P.H. (1995). Techniques for verifying the accuracy of risk measurement models. *Journal of Derivatives*, 3(2), pp. 73–84.

[2] Christoffersen, P. (1998). Evaluating interval forecasts. *International Economic Review*, 39(4), pp. 841–862.

[3] Markowitz, H. (1952). Portfolio selection. *Journal of Finance*, 7(1), pp. 77–91.

[4] Morgan, J.P. and Reuters (1996). *RiskMetrics Technical Document*, 4th edition. J.P. Morgan, New York.

[5] Jorion, P. (2007). *Value at Risk: The New Benchmark for Managing Financial Risk*, 3rd edition. McGraw-Hill, New York.

[6] McNeil, A.J., Frey, R. and Embrechts, P. (2005). *Quantitative Risk Management: Concepts, Techniques and Tools*. Princeton University Press, Princeton, NJ.

[7] Basel Committee on Banking Supervision (2004). *International Convergence of Capital Measurement and Capital Standards* (Basel II). Bank for International Settlements, Basel.